


Rich Feature
Representation


Learning Rich Features from RGB-D Images for Object Detection and Segmentation



(a) Image classification



(b) Object localization



(c) Semantic segmentation

Agenda →

INTRODUCTION

01

Intro to the Research
paper we will be
presenting.

Agenda →

01

INTRODUCTION

Intro to the Research paper we will be presenting.

02

OVERVIEW OF THE PROBLEM

What problem is this research paper attempting to solve?

Why this problem is interesting and what challenges it poses?

Agenda →

01

INTRODUCTION

Intro to the Research paper we will be presenting.

02

OVERVIEW OF THE PROBLEM

What problem is this research paper attempting to solve?

Why this problem is interesting and what challenges it poses?

03

SCOPE OF THE PROJECT

Defining the exact scope of what we are attempting as a part of the project.

What are the specific goals and objectives of our work?

Agenda →

01

INTRODUCTION

Intro to the Research paper we will be presenting.

02

OVERVIEW OF THE PROBLEM

What problem is this research paper attempting to solve?

Why this problem is interesting and what challenges it poses?

03

SCOPE OF THE PROJECT

Defining the exact scope of what we are attempting as a part of the project.

What are the specific goals and objectives of our work?

04

METHODOLOGY

How do we plan to accomplish our goals?

The dataset we will be using and how to access it.
The baseline model we will be using.
Experiments, ablations, or comparisons we plan to conduct.
Computing resources we will be using
and other details ...

Agenda →

TIMELINE

Outlining a timeline for how
we plan to complete the
above tasks

Identification of any major
milestones or deadlines

05

Agenda →

TIMELINE

05

Outlining a timeline for how we plan to complete the above tasks

Identification of any major milestones or deadlines

06

INITIAL PROGRESS

Initial progress we have made on the code level

Explain any challenges we have faced and how we have overcome them

Agenda →

TIMELINE

05

Outlining a timeline for how we plan to complete the above tasks

Identification of any major milestones or deadlines

INITIAL PROGRESS

06

Initial progress we have made on the code level

Explain any challenges we have faced and how we have overcome them

CONCLUSION

07

Summarize the key points of the presentation

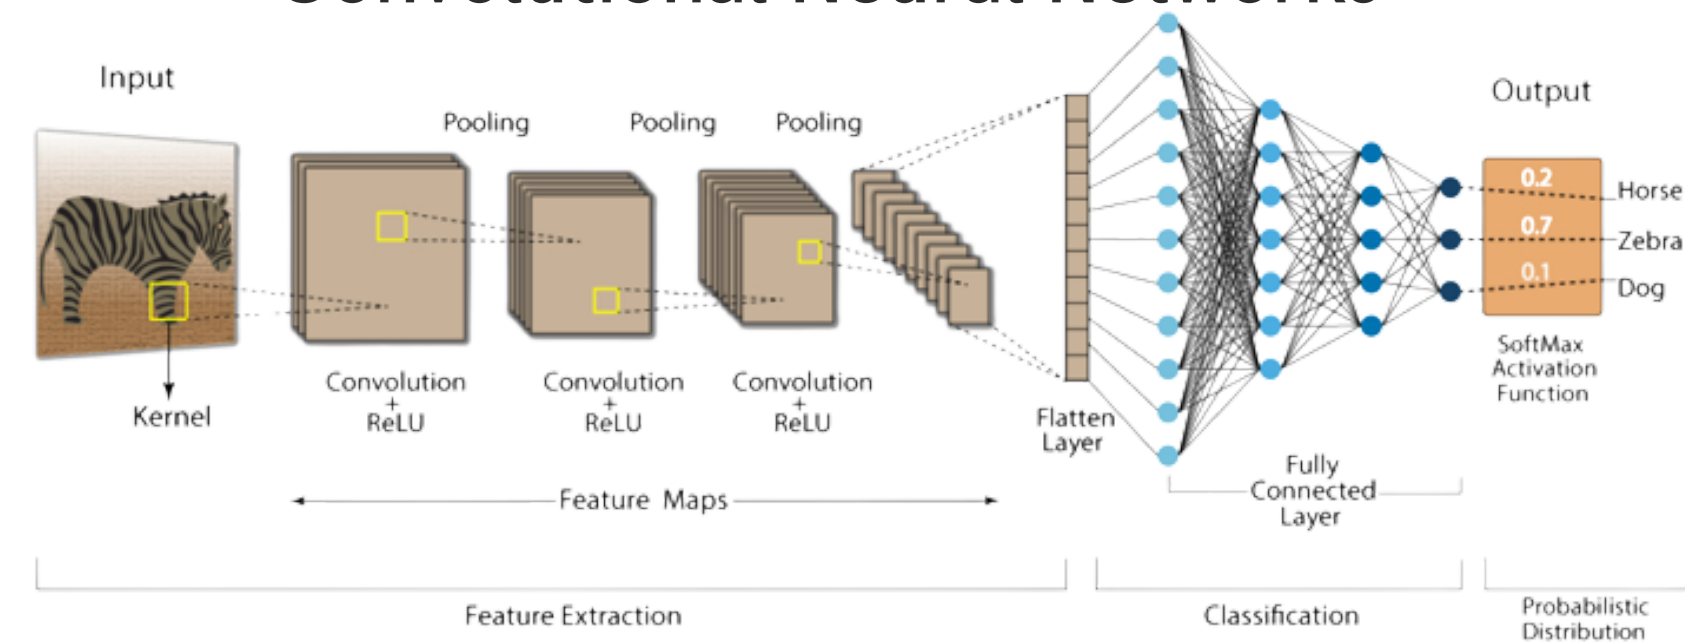
Explain the potential impact of our work

Invite questions and discussion from the audience


Introduction to Problem Statement



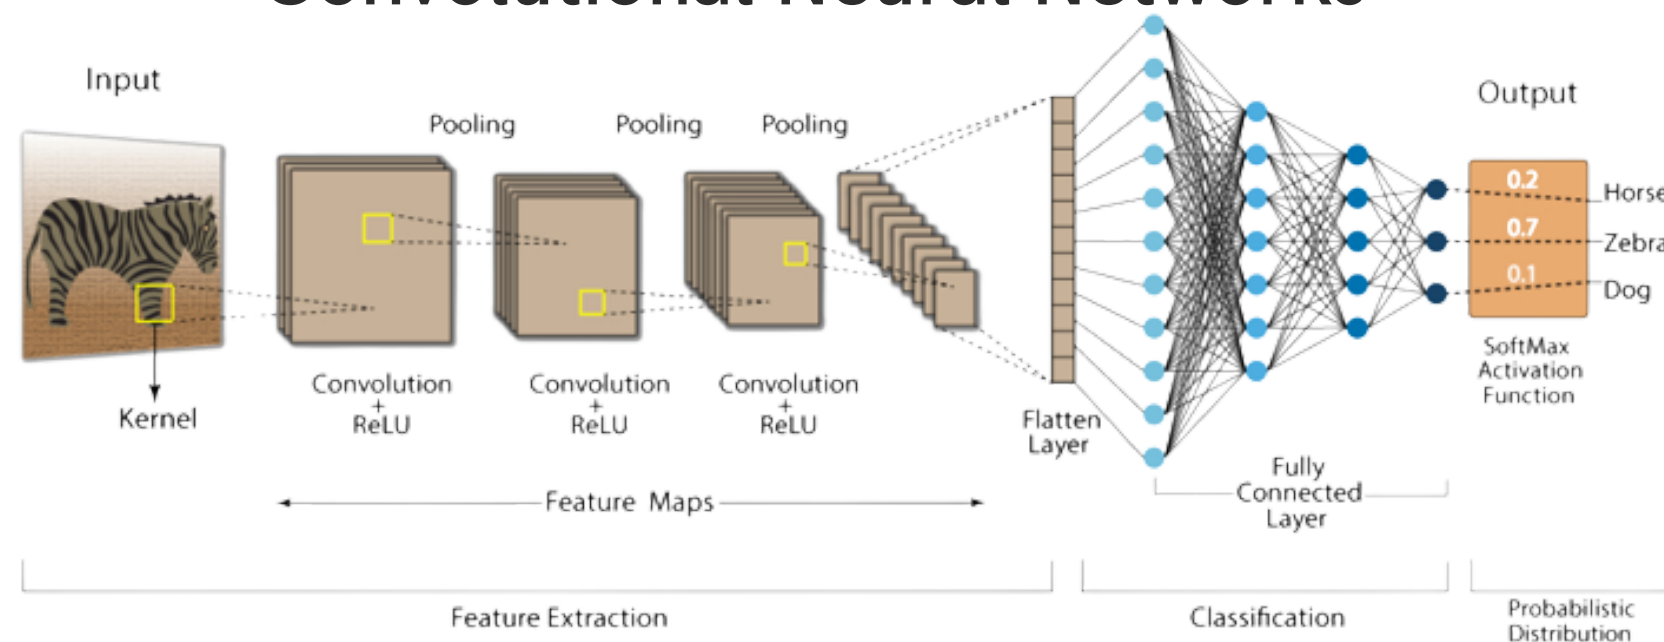
Strengthening Vision Perception via Convolutional Neural Networks



Introduction to Problem Statement




Strengthening Vision Perception via Convolutional Neural Networks



Learning Rich Feature Representations!



(a) Image classification

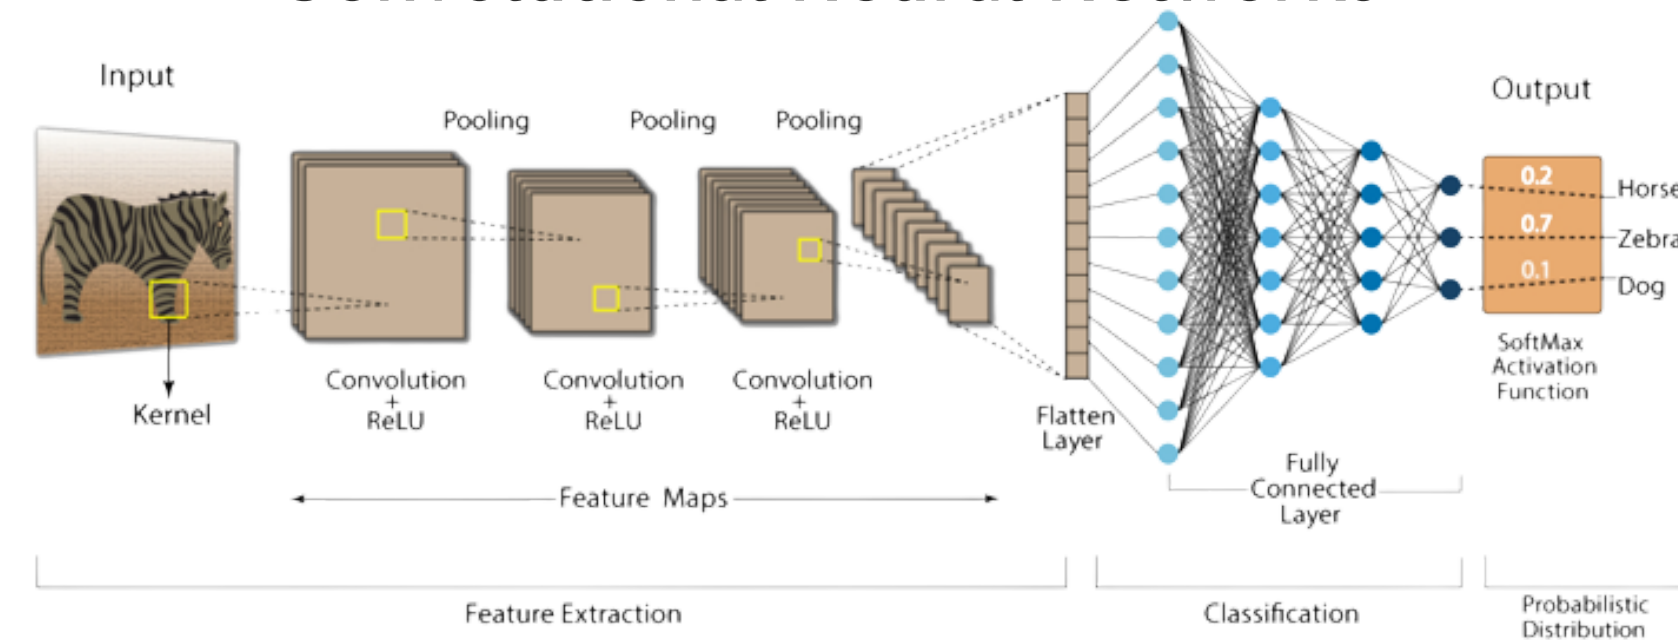


(b) Object localization

Introduction to Problem Statement




Strengthening Vision Perception via Convolutional Neural Networks



Learning Rich Feature Representations!





(a) Image classification



(b) Object localization


Easy for human cognition - difficult for machines



Introduction to Problem Statement




Strengthening Vision Perception via Convolutional Neural Networks



Learning Rich Feature Representations!





(a) Image classification



(b) Object localization

Easy for human cognition - difficult for machines

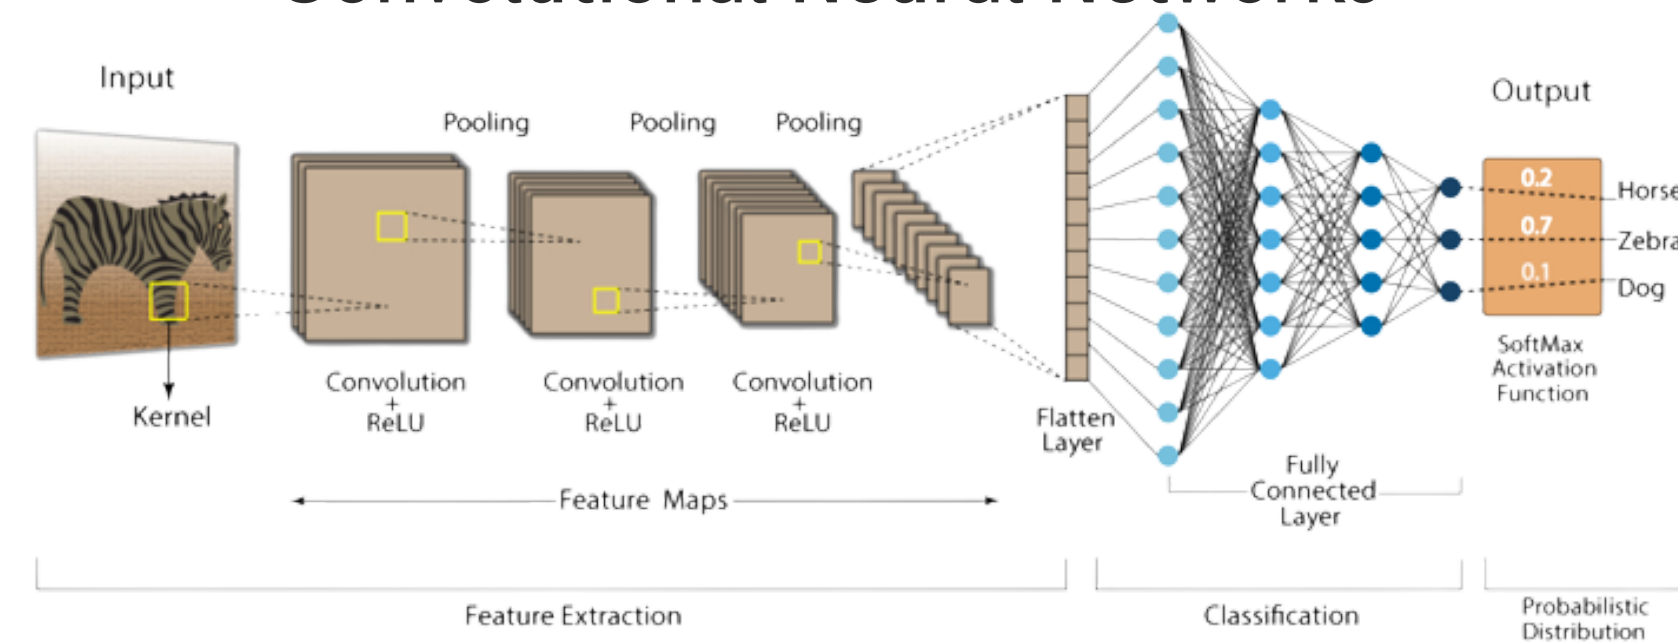


(c) Semantic segmentation

Introduction to Problem Statement




Strengthening Vision Perception via Convolutional Neural Networks



Learning Rich Feature Representations!





(a) Image classification




(b) Object localization

Easy for human cognition - difficult for machines



(c) Semantic segmentation



(d) Instance segmentation

Overview: Research Idea →

- Develop an Integrated system for scene understanding from **RGB-D** images




Overview: Research Idea →

- Develop an Integrated system for scene understanding from **R****G****B**-D images

Object Detection

- Geocentric embedding instead of raw depth images to CNN = better feature representation (56% relative improvement).



Overview: Research Idea →

- Develop an Integrated system for scene understanding from **RGB-D** images

- Decision forest approach that labels pixels of object instances as foreground or background by using the geocentric pose features.

Object Detection

- Geocentric embedding instead of raw depth images to CNN = better feature representation (56% relative improvement).

Instance Segmentation



Overview: Research Idea →

- Develop an Integrated system for scene understanding from RGB-D images

Object Detection

- Geocentric embedding instead of raw depth images to CNN = better feature representation (56% relative improvement).

- Decision forest approach that labels pixels of object instances as foreground or background by using the geocentric pose features.

Instance Segmentation

Semantic Segmentation

- Output of object detectors in a supapixel classification framework for semantic scene segmentation (24% relative improvement).



Overview: Proposed Pipeline →

Architecture of the integrated system

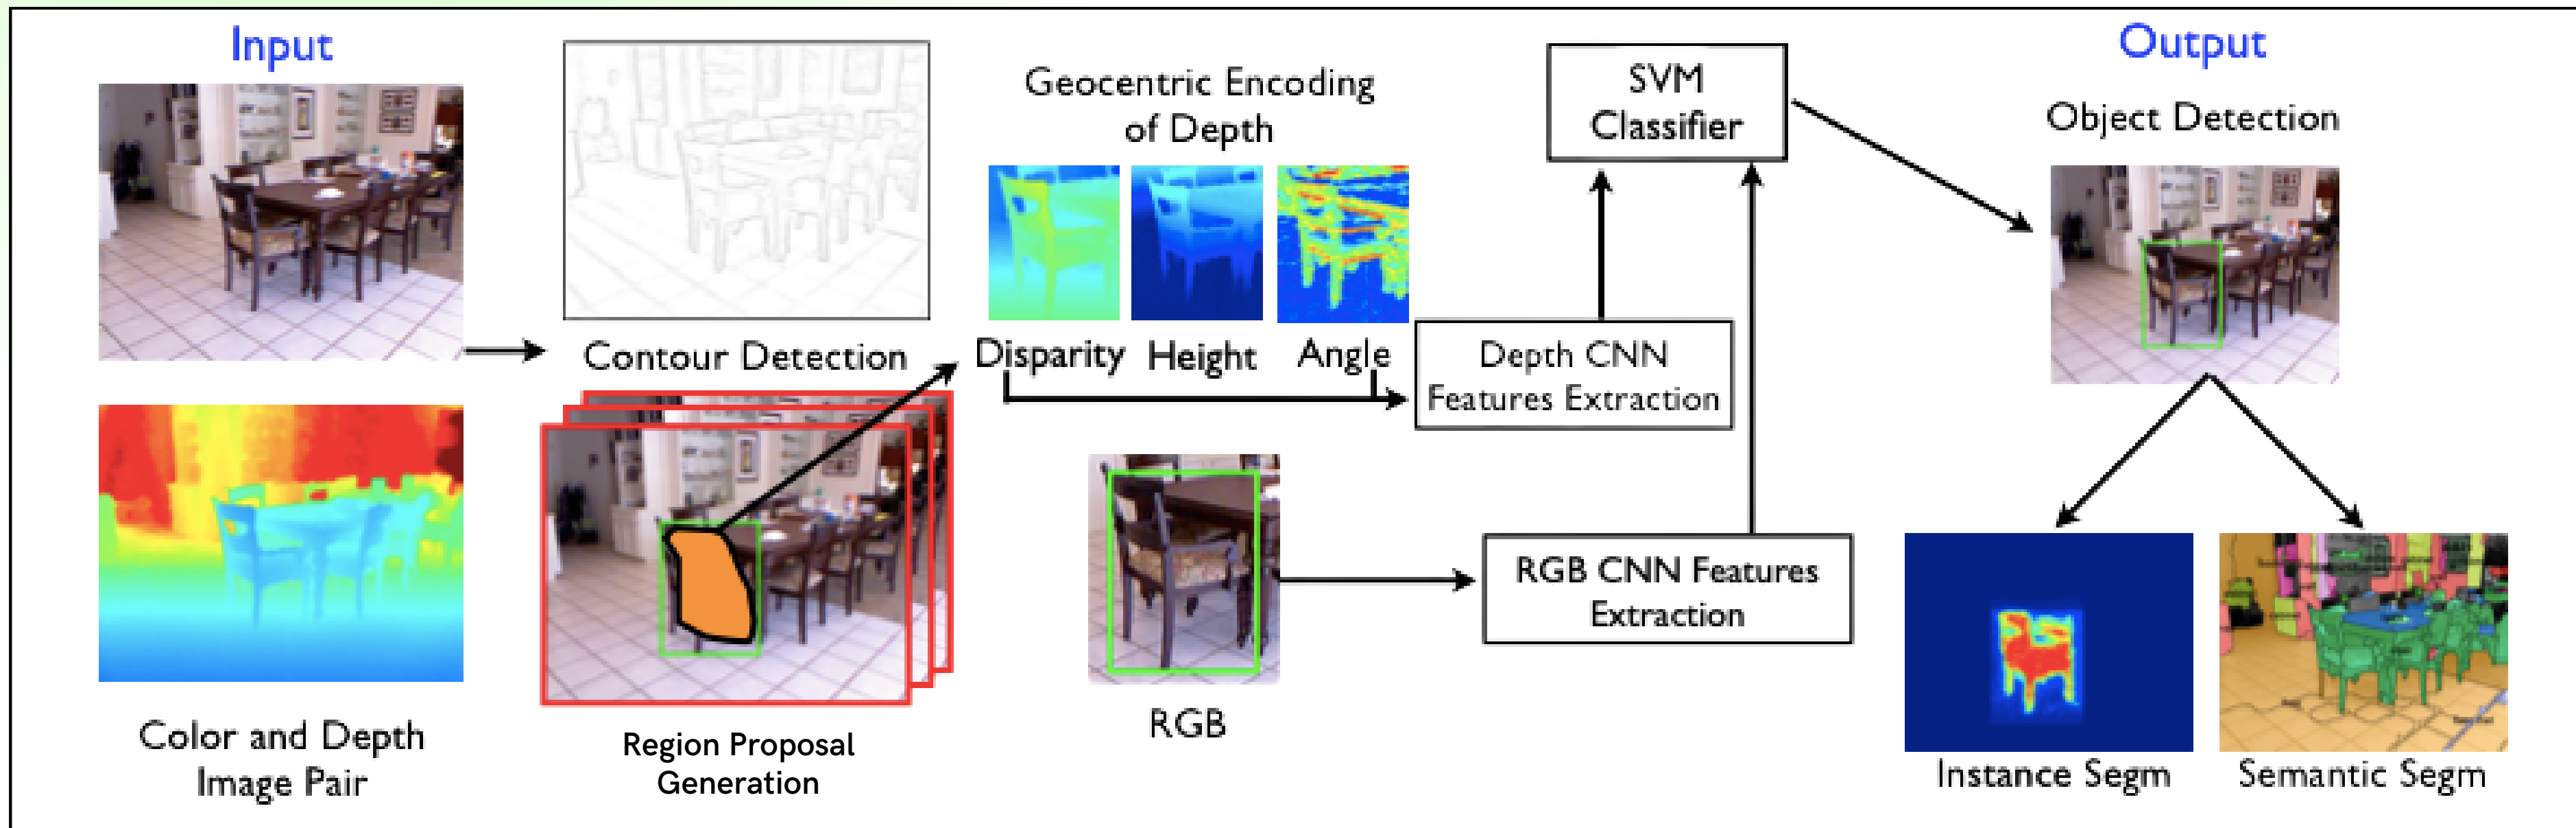
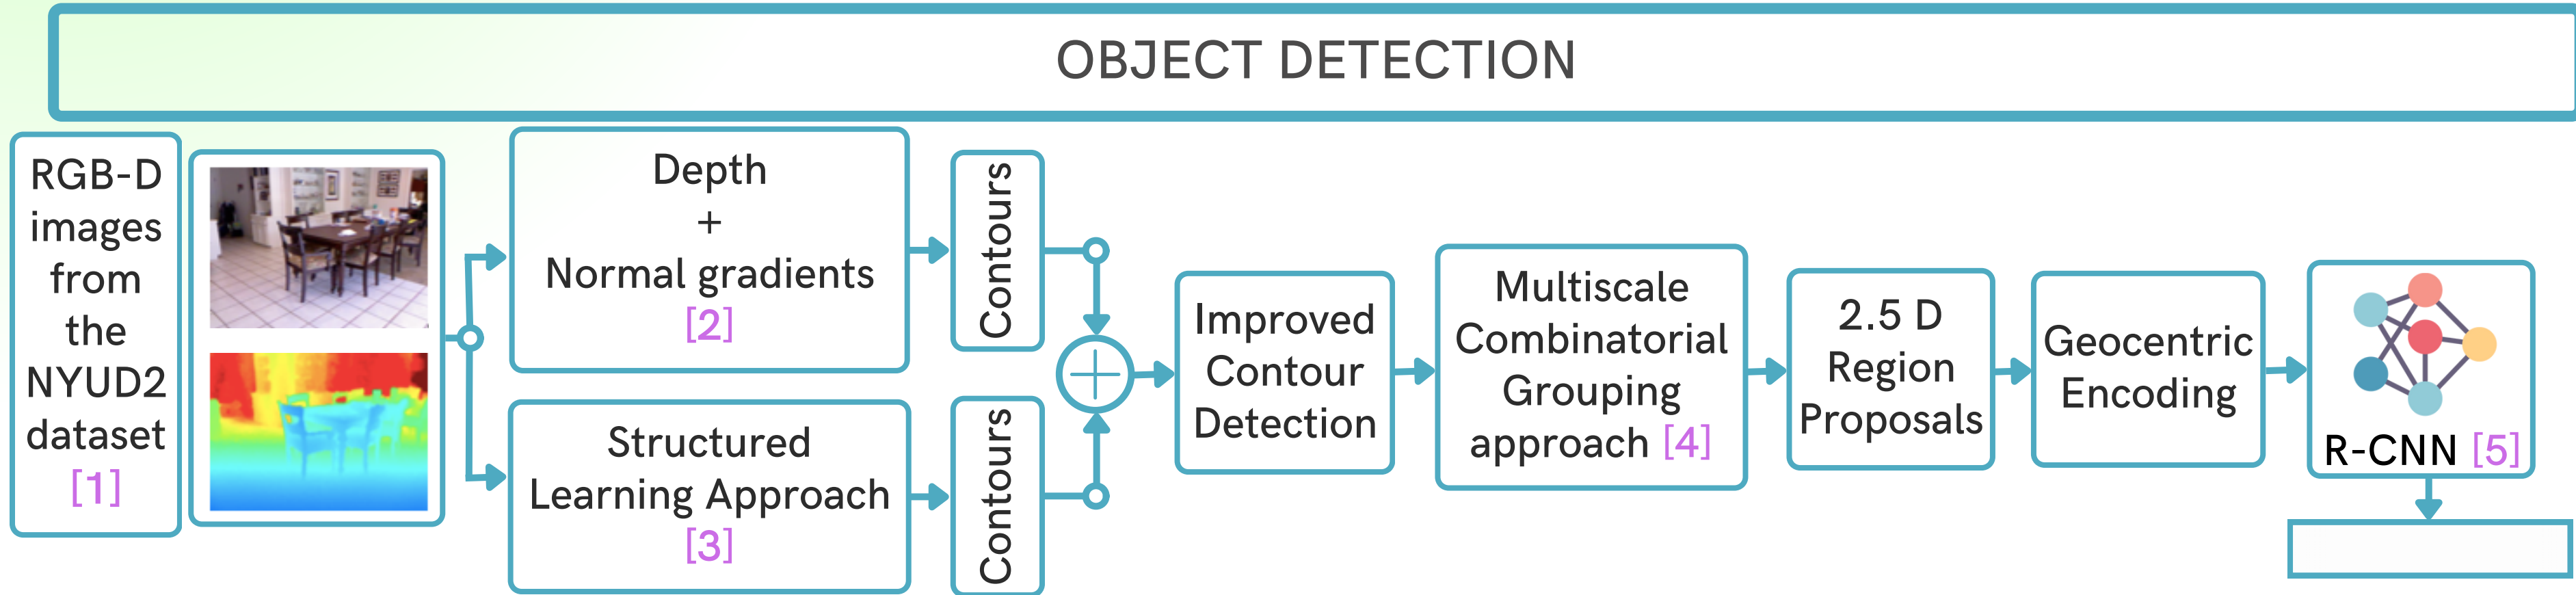
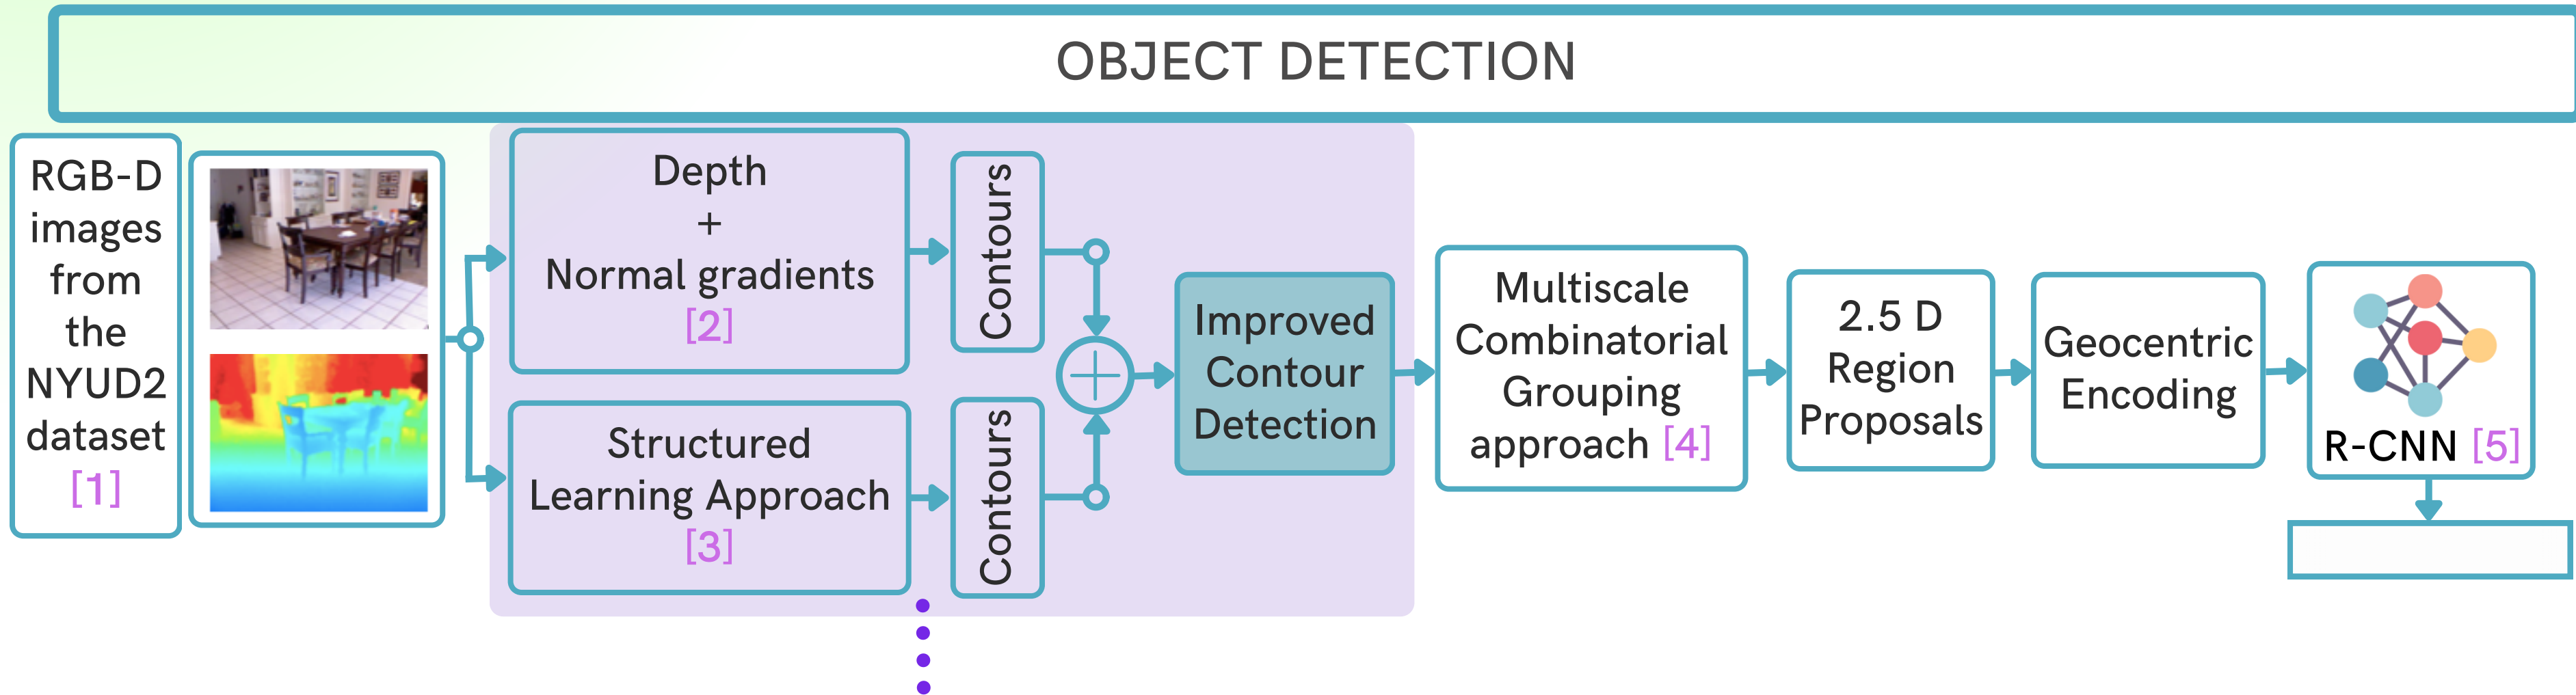


Fig 1. Overview: from an RGB and depth image pair, the system detects contours, generates 2.5D region proposals, classifies them into object categories, and then infers segmentation masks for instances of “thing”-like objects, as well as labels for pixels belonging to “stuff”-like categories.

Architecture of the integrated system



Architecture of the integrated system



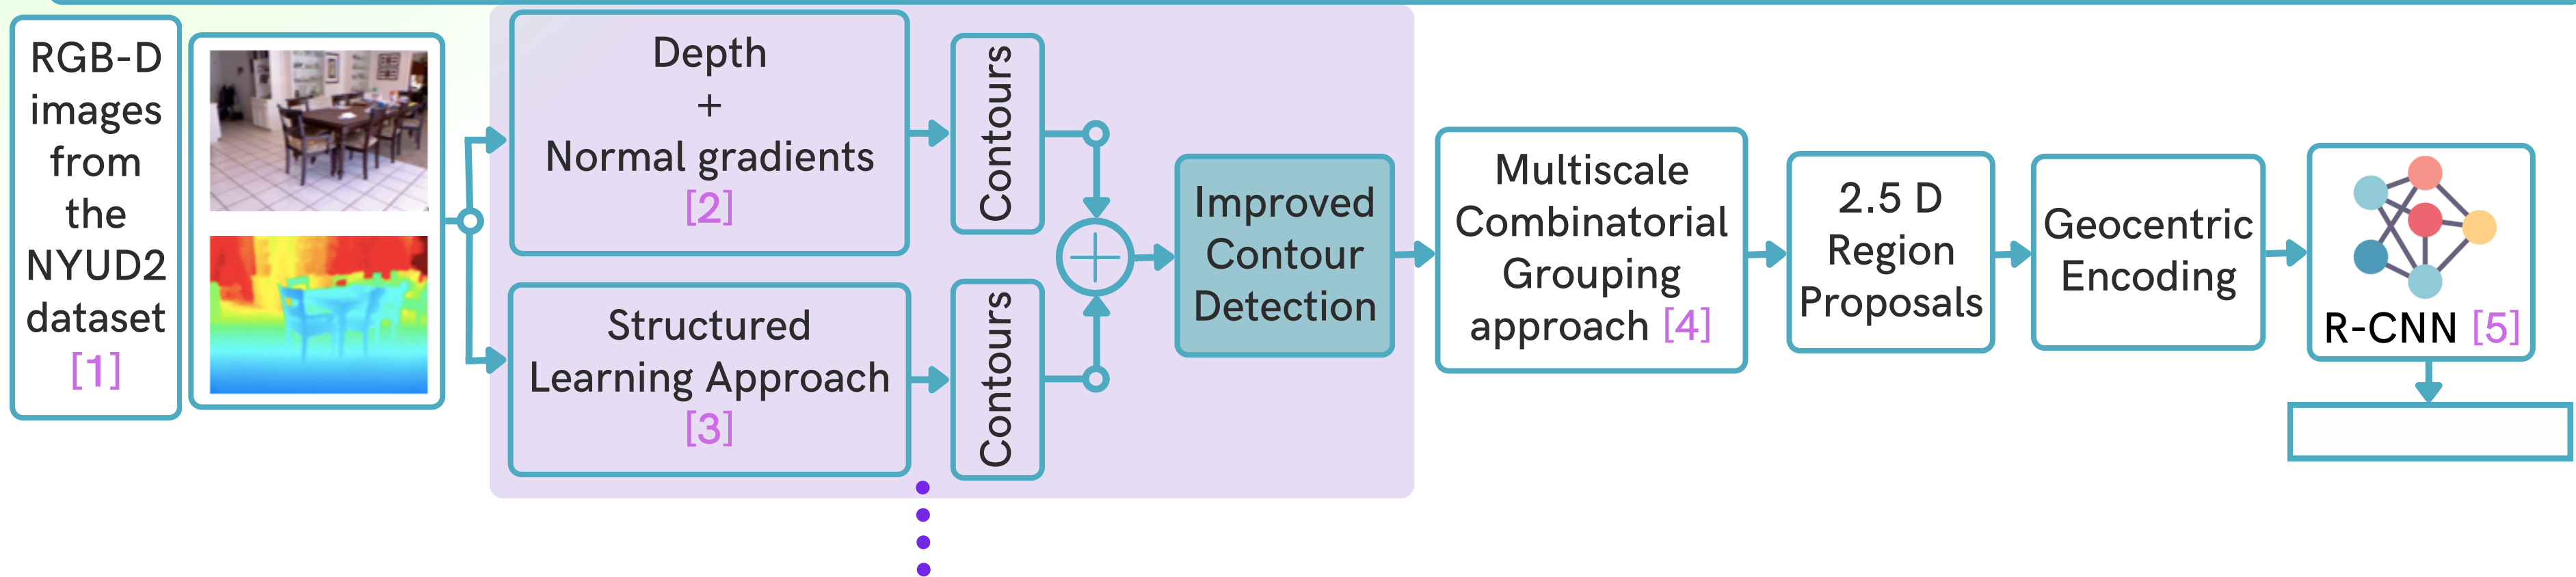
[2] uses gPb_ucm [6] and proposed local geometric gradients dubbed NG-, NG+ and DG to capture convex, concave normal gradients and depth gradients

[3] proposed an approach based on structured random forests to directly classify a pixel as being a contour pixel or not

[3] produces better localized contours that capture fine details, but tend to miss normal discontinuities that [2] easily finds

Architecture of the integrated system

OBJECT DETECTION




[2] uses gPb_ucm [6] and proposed local geometric gradients dubbed NG-, NG+ and DG to capture convex, concave normal gradients and depth gradients

[3] proposed an approach based on structured random forests to directly classify a pixel as being a contour pixel or not

[3] produces better localized contours that capture fine details, but tend to miss normal discontinuities that [2] easily finds

[2]

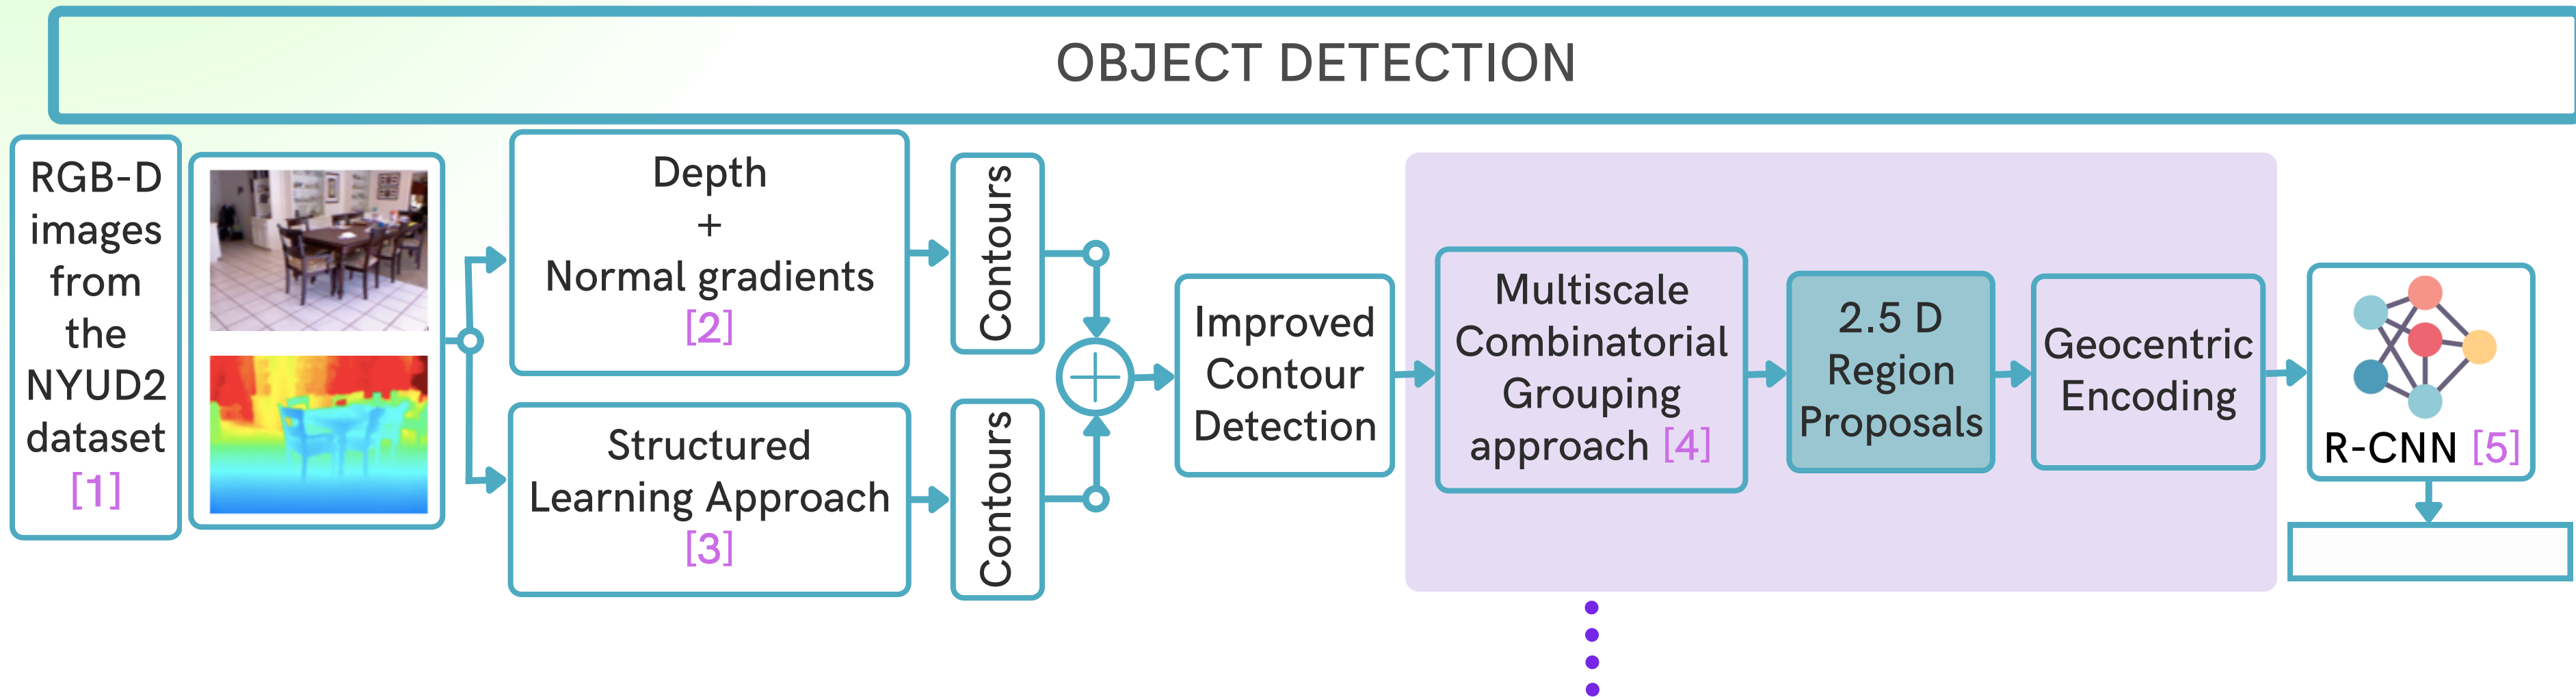


[3]




Contours from proposed approach

Architecture of the integrated system




MCG is an approach used for object proposal generation.

From improved contour, we obtain object proposals by generalizing MCG to RGB-D images



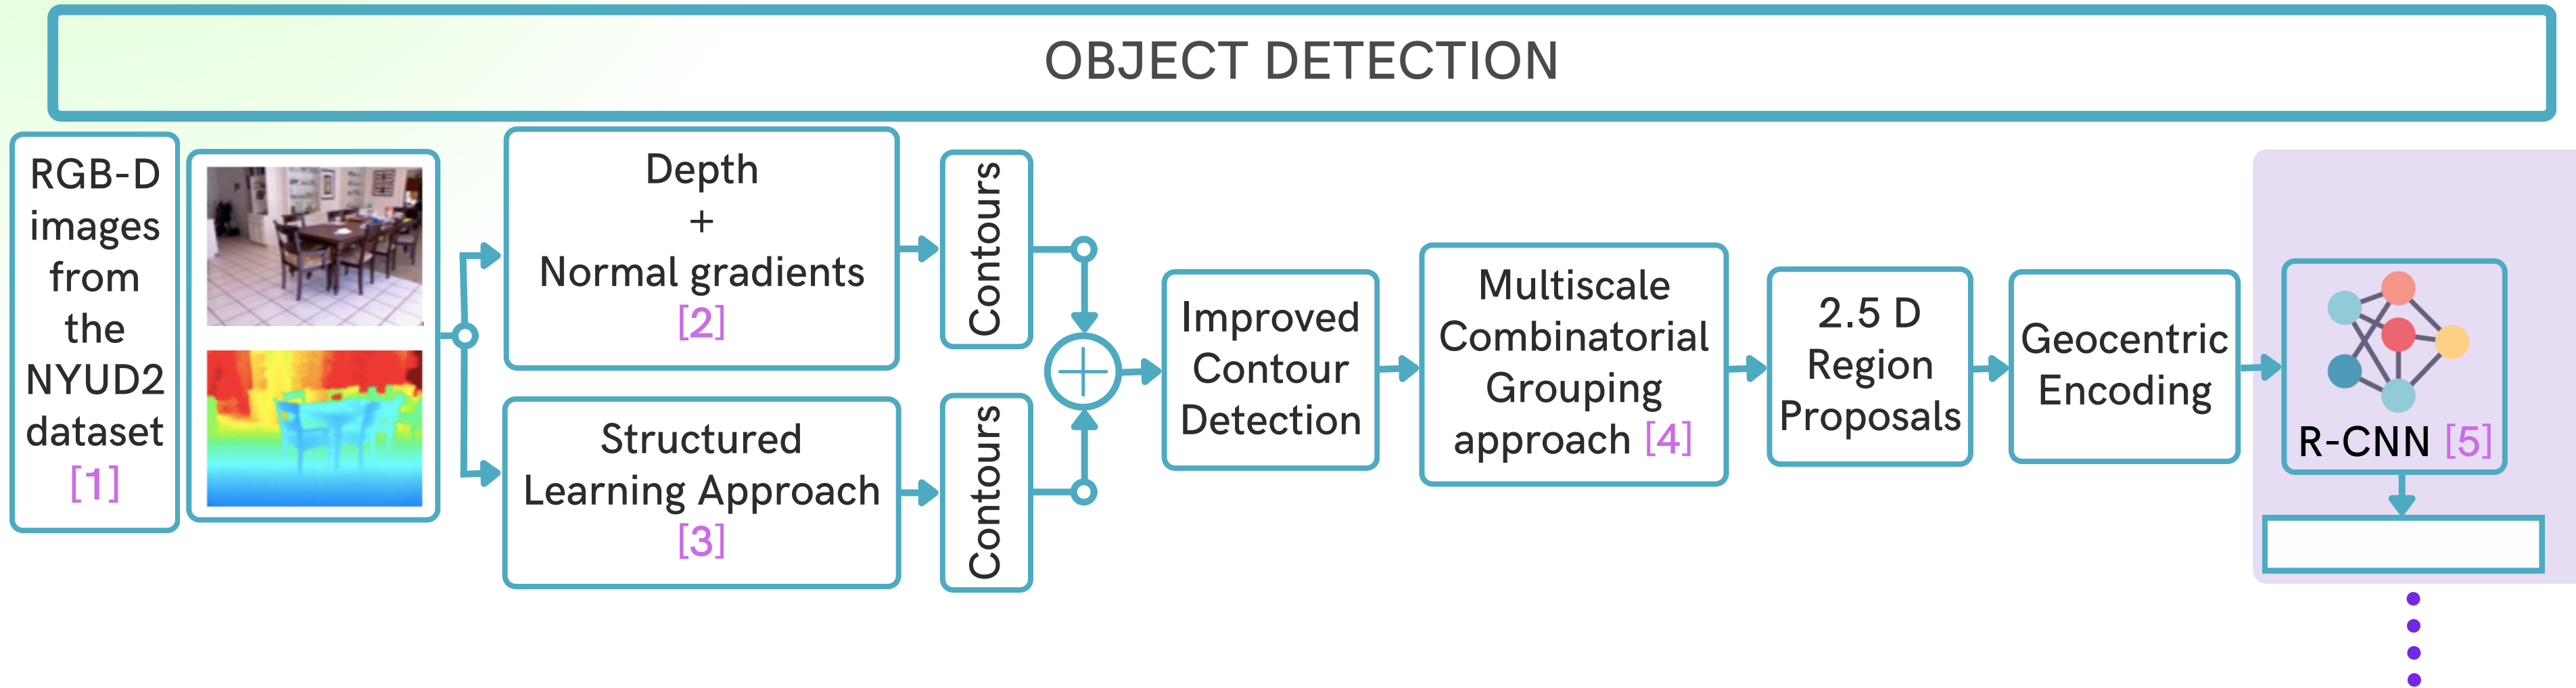
Geocentric Encoding of Depth



Disparity Height Angle

Geocentric encoding allows CNN to learn strong features than by using disparity alone

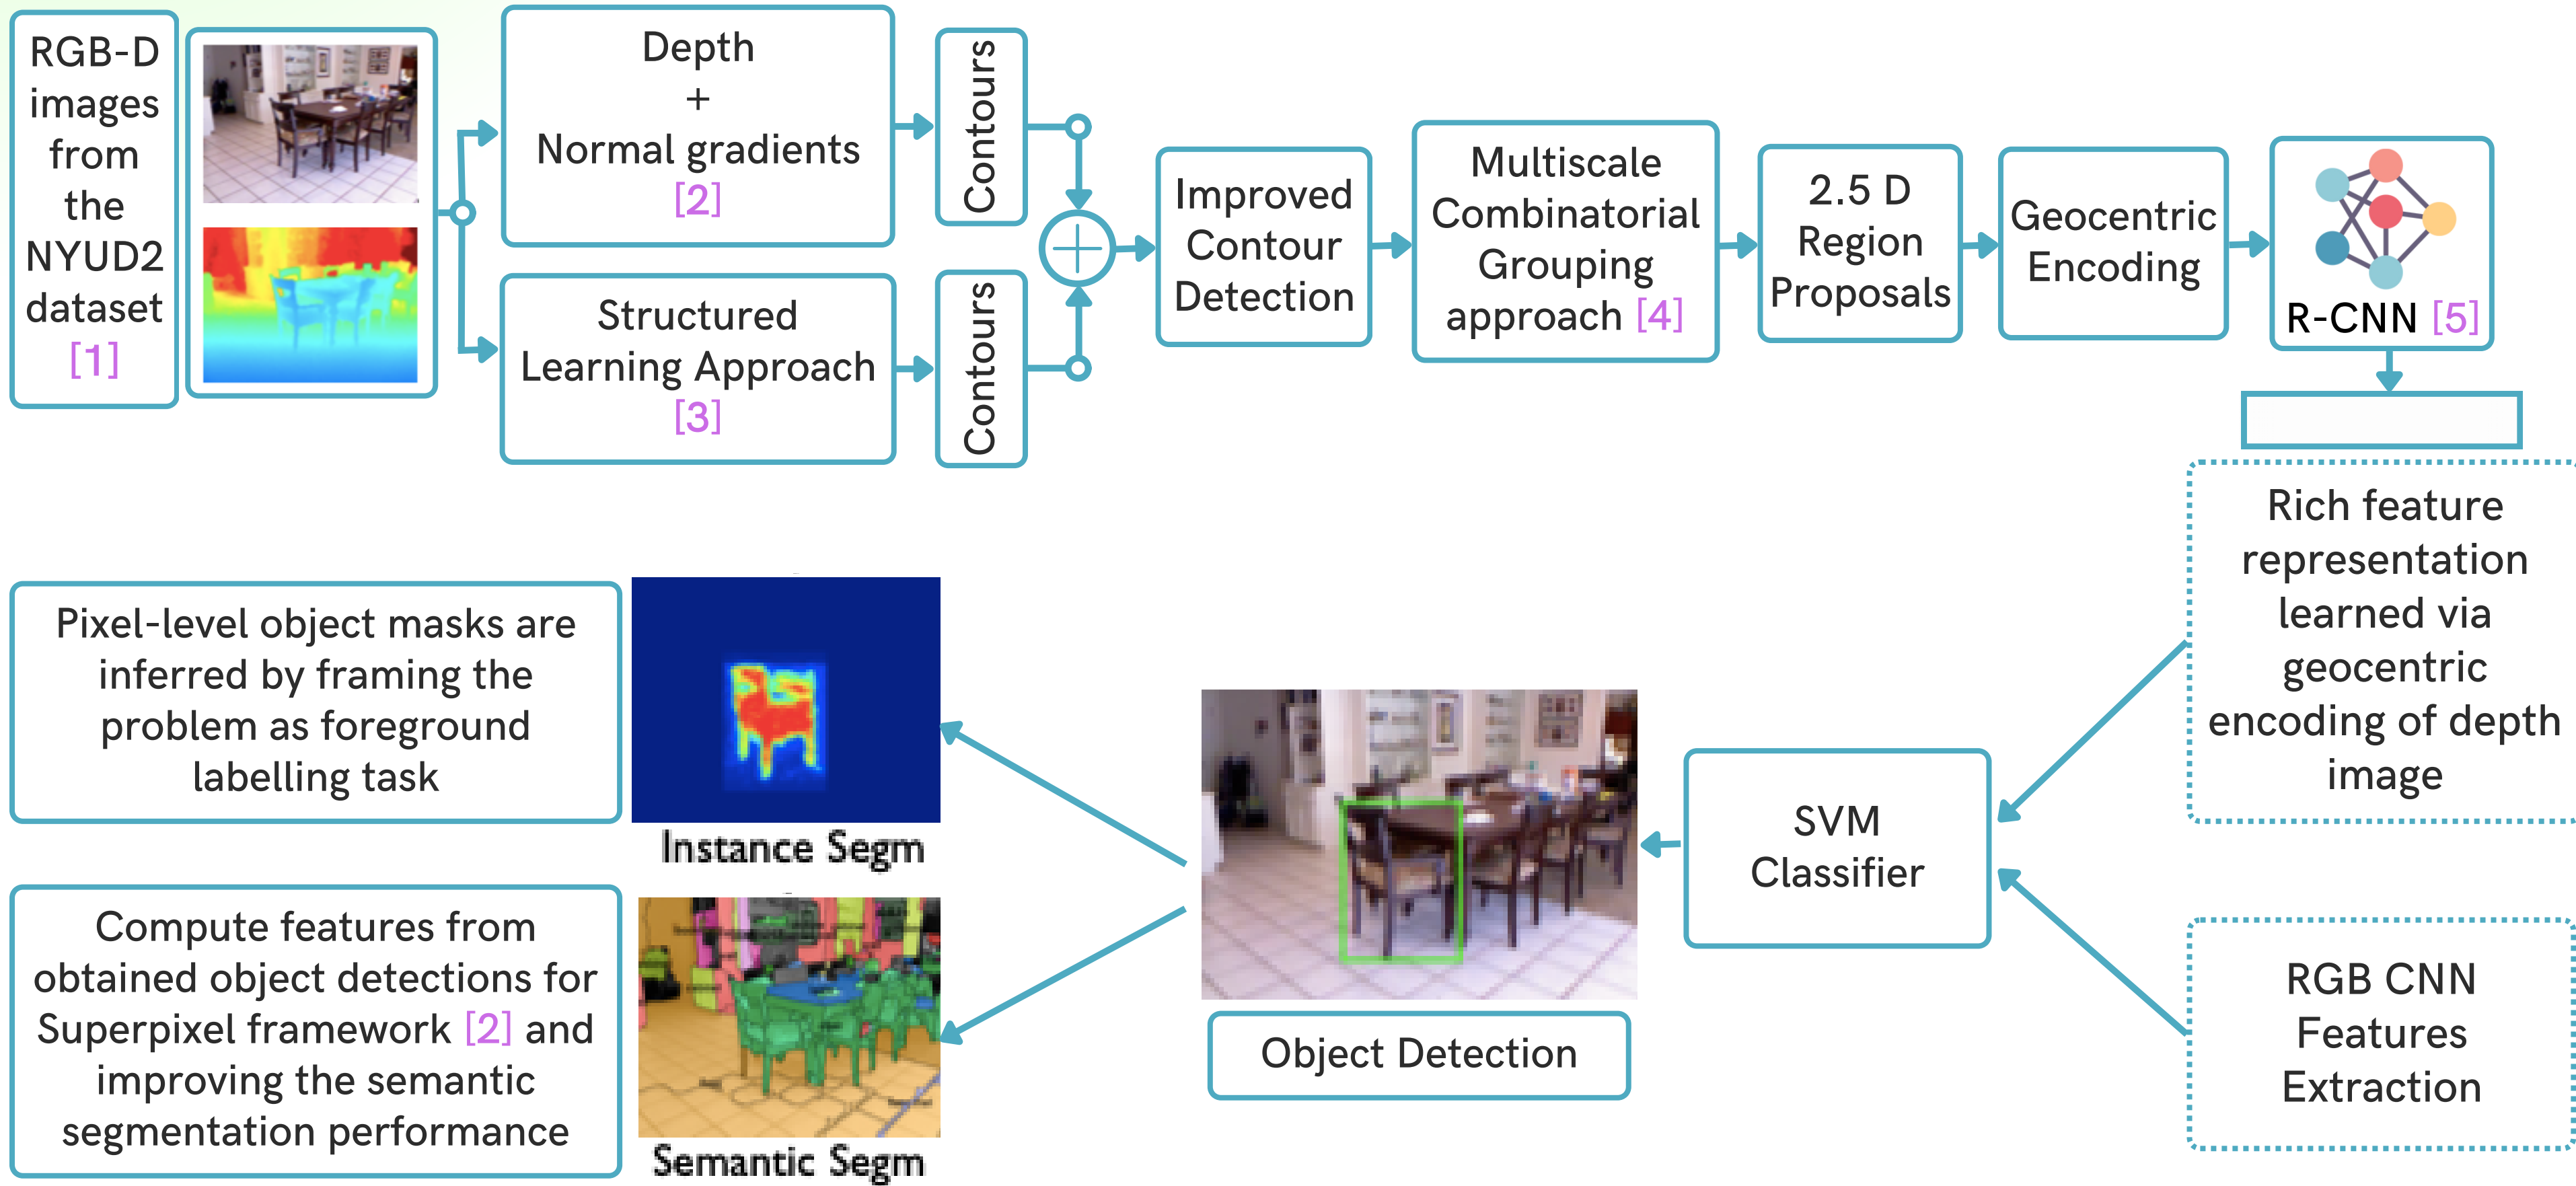
Architecture of the integrated system



A large R-CNN pre-trained on RGB images, can be adapted to generate rich features from this representation of depth images.

Architecture of the integrated system

OBJECT DETECTION + INSTANCE SEGMENTATION + SEMANTIC SEGMENTATION



RESULTS




Fig 2: Precision-Recall curve on boundaries on the NYUD2 dataset

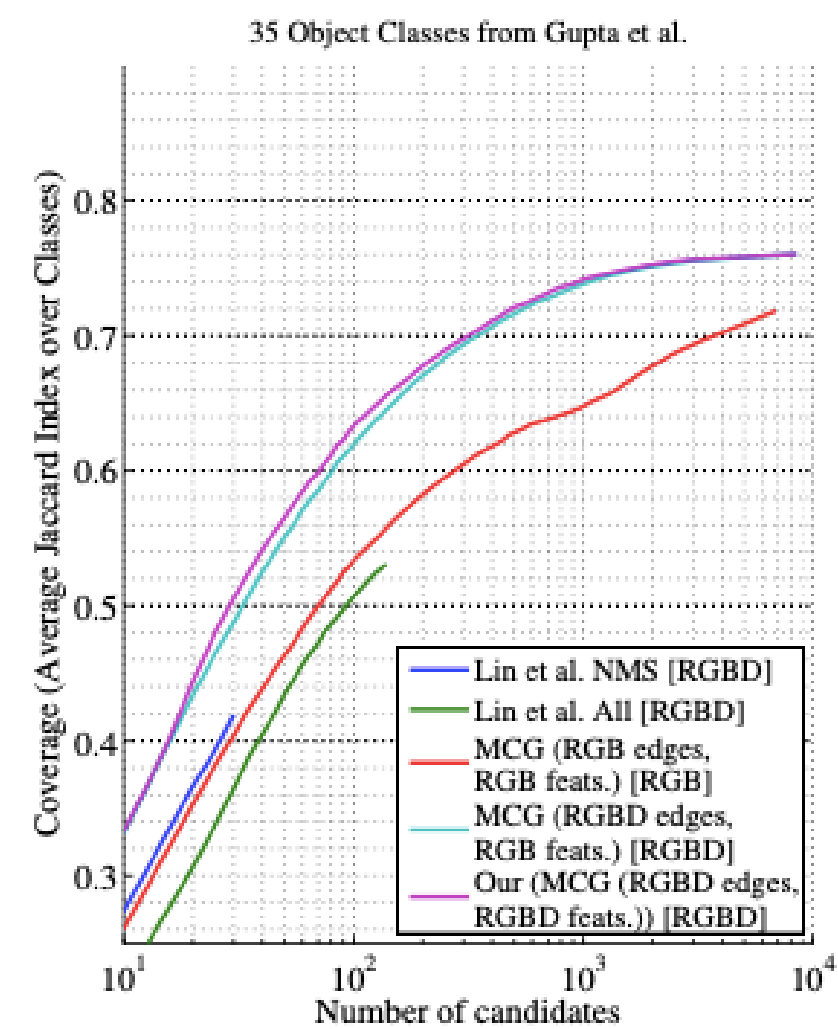


Fig. 4. Region Proposal Quality: Coverage as a function of the number of region proposals per image: Depth-based region proposals using the improved RGB-D contours work better than all past approaches, while at the same time being more general. Note that the X-axis is on a log scale.

OBJECT DETECTION

Final object detection system achieves an average precision of 37.3%, which is a 56% relative improvement over existing methods

Why is this Problem Interesting




FACILITATE THE USE OF PERCEPTION IN FIELDS LIKE SMART TRAFFIC MONITORING, SUPPLY-CHAIN MONITORING, MEDICAL ASSISTANCE

- Advanced human-machine interfaces (gesture recognition, motion tracking, etc.)
- Crowd monitoring
- Self-driving cars
- Navigation systems for unmanned robotic vehicles, and much more



Virtual Reality



Robot Navigation



Augmented Reality

Challenges →

Depth images may contain noise due to sensor limitations or environmental factors such as reflections or interference from other devices. This can lead to inaccurate depth measurements, which can affect the accuracy of object detection and segmentation

Processing both RGB and depth information requires more computational resources than traditional 2D image-based methods. This can lead to slower detection and segmentation times, which may be a problem in real-time applications

Objects can vary in shape, size, color, and texture, which can make it difficult to create accurate models for object detection and segmentation. This can lead to false positives or false negatives

Training deep learning models for RGB-D object detection and instance segmentation requires large amounts of annotated data. However, there are currently only a limited number of publicly available annotated datasets for this task



Data Quality



Computational Complexity



Variability in object appearance

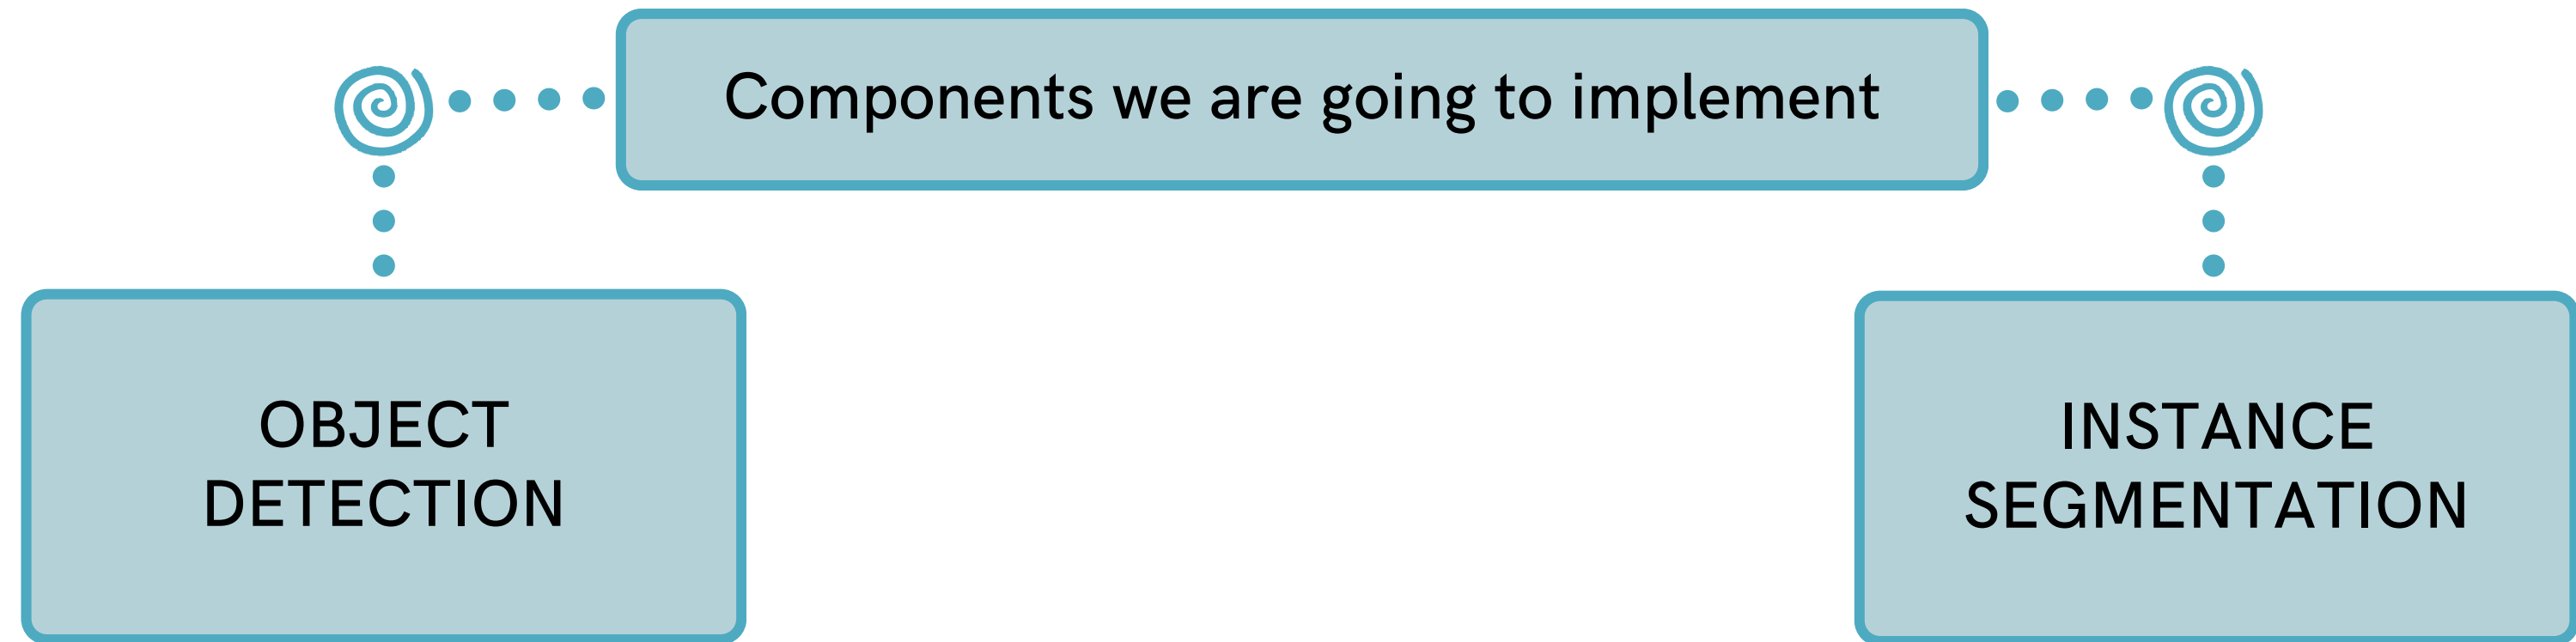


Limited availability of annotated dataset

Scope of the Project



What are we attempting as a part of the project



Exact details regarding the implementation of all the sub-components - Timeline Slide

Methodology →

1

Datasets

What and how to access?

Dataset name

NYU Depth
Dataset v2

Access

Open Source

2

Computing Resources

What and how much?

Specifications

GPU: RTX
2080Ti
CPU: Intel Xeon
E5-2640 v4

Estimated Time

Finetuning: 90
hrs (100 epochs)

3

Baseline models

The requirements

2.5D region
proposal

Geocentric
encoding of
depth
information
(HHA)

Random forest
for **Instance
Segmentation**

Finetune
R-CNN for
HHA

Implement
**Semantic
Segmentation**

4

Improvements

The extras

Replace R-
CNN with
faster R-CNN

Using v2 of
NYUD dataset

Replace SVM
classifier with a
Deep Learning
base method

Timeline →



4 Apr

Write code for 2.5 D region proposal.



8 Apr

Implement geocentric encoding of depth information



10 Apr

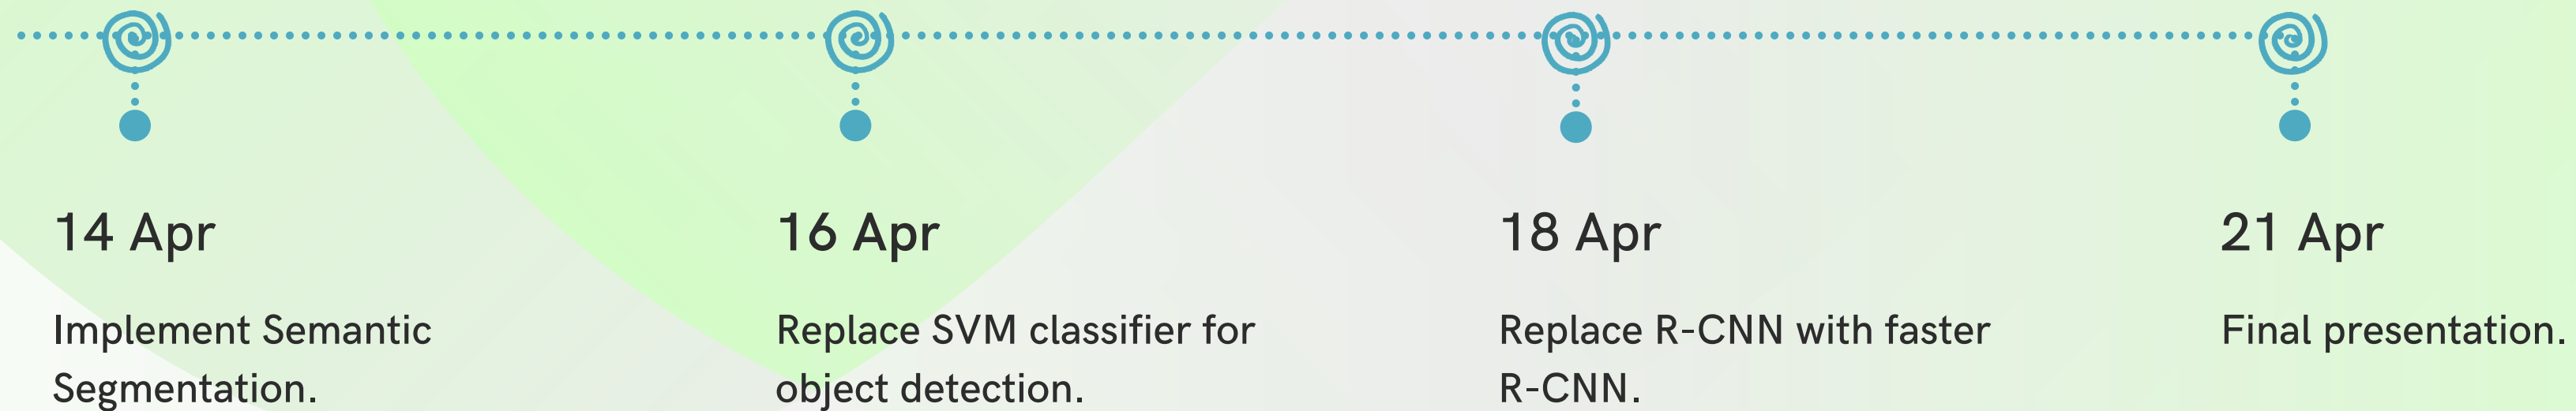
Fine tune R-CNN for geocentric encoding.



12 Apr

Random forest for instance segmentation.


Timeline →



Initial Progress →

Implementation of 2.5D region proposal

[4] Arbeláez, P., Pont-Tuset, J., Barron, J., Marques, F., Malik, J.: Multiscale combinatorial grouping. In: CVPR (2014)




Region Proposal
Generation

Encoding Depth Images for Feature Learning

[18] Gupta, S., Arbeláez, P., Malik, J.: Perceptual organization and recognition of indoor scenes from RGB-D images. In: CVPR (2013)

Geocentric Encoding of Depth



Disparity Height Angle

Team →



Deepti Rawat

deepti.rawat@research.iiit.ac.in

Roll No: 2022801016



Rishabh Khanna

rishabh.khanna@research.iiit.ac.in

Roll No: 2019113025



Mihir Rawat

mihir.r@research.iiit.ac.in

Roll No: 2020113005



Kshitijaa Jaglan

kshitijaa.jaglan@research.iiit.ac.in

Roll No: 2019115005

Work Distribution →

Feature Representation



2.5D Region Proposal

Kshitijaa, Mihir



Geocentric Encoding

Deepti, Rishabh



Fine Tuning R-CNN

Deepti, Rishabh



RGB R-CNN

Mihir, Kshitijaa

Segmentation Pipeline



SVM Classifier

Deepti, Mihir



Instance Segmentation

Rishabh, Kshitijaa



Semantic Segmentation

Deepti, Mihir

Proposed Changes



Faster R-CNN

Mihir, Kshitijaa



DL Classifier

Kshitijaa



Instance Segmentation

Deepti, Rishabh

References



- [1] https://cs.nyu.edu/~silberman/datasets/nyu_depth_v2.html
- [2] Gupta, S., Arbela'ez, P., Malik, J.: Perceptual organization and recognition of indoor scenes from RGB-D images. In: CVPR (2013)
- [3] Dolla'r, P., Zitnick, C.L.: Structured forests for fast edge detection. In: ICCV (2013)
- [4] Arbela'ez, P., Pont-Tuset, J., Barron, J., Marques, F., Malik, J.: Multiscale combinatorial grouping. In: CVPR (2014)
- [5] Girshick, R., Donahue, J., Darrell, T., Malik, J.: Rich feature hierarchies for accurate object detection and semantic segmentation. In: CVPR (2014)
- [6] Arbela'ez, P., Maire, M., Fowlkes, C., Malik, J.: Contour detection and hierarchical image segmentation. TPAMI (2011)